

Predicting species richness and diversity

Bert van der Veen

University of Bayreuth

Questions so far?



Diversity as a community property

A JSDM produces a lot of output: intercepts, slopes, associations, predicted probabilities for every species at every site. Rich, but hard to communicate.

Diversity as a community property

A JSDM produces a lot of output: intercepts, slopes, associations, predicted probabilities for every species at every site. Rich, but hard to communicate.

Diversity indices — species richness, Shannon entropy, Simpson's index — are scalar summaries of the community matrix $\mathbf{Y}$.

Diversity as a community property

A JSDM produces a lot of output: intercepts, slopes, associations, predicted probabilities for every species at every site. Rich, but hard to communicate.

Diversity indices — species richness, Shannon entropy, Simpson's index — are scalar summaries of the community matrix $\mathbf{Y}$.

If we model the community, diversity follows.

Why model-based diversity?

Indices such as Shannon entropy and Simpson's index have no generative model: there is no likelihood, no data-generating process, nothing to fit or check.

Why model-based diversity?

Indices such as Shannon entropy and Simpson's index have no generative model: there is no likelihood, no data-generating process, nothing to fit or check.

Without a generative model:

- ▶ No uncertainty quantification: no standard errors, no confidence intervals
- ▶ No ability to account for sampling: effort, detection probability, protocol differences
- ▶ No ability to account for data properties: overdispersion, zero-inflation, censoring
- ▶ No prediction for new sites, times, or conditions

Why model-based diversity?

Indices such as Shannon entropy and Simpson's index have no generative model: there is no likelihood, no data-generating process, nothing to fit or check.

Without a generative model:

- ▶ No uncertainty quantification: no standard errors, no confidence intervals
- ▶ No ability to account for sampling: effort, detection probability, protocol differences
- ▶ No ability to account for data properties: overdispersion, zero-inflation, censoring
- ▶ No prediction for new sites, times, or conditions

A community model addresses all of these: diversity follows from

What is species richness?

Species richness SR_i at site i is the number of species recorded:

$$SR_i = \sum_{j=1}^m y_{ij} \quad (\text{for binary data}) \quad (1)$$

What is species richness?

Species richness SR_i at site i is the number of species recorded:

$$SR_i = \sum_{j=1}^m y_{ij} \quad (\text{for binary data}) \quad (1)$$

It is by far the most common biodiversity metric: simple to calculate, intuitive to interpret, and widely used in conservation and ecology.

What is species richness?

Species richness SR_i at site i is the number of species recorded:

$$SR_i = \sum_{j=1}^m y_{ij} \quad (\text{for binary data}) \quad (1)$$

It is by far the most common biodiversity metric: simple to calculate, intuitive to interpret, and widely used in conservation and ecology.

But it is a **derived quantity** from multivariate community data. Analyzing it directly means throwing away most of the data.

The typical approach

Species richness is most commonly analyzed with Poisson or negative-binomial regression:

$$\log\{\mathbb{E}(SR_i)\} = \alpha + \mathbf{x}_i^\top \boldsymbol{\beta} \quad (2)$$

The typical approach

Species richness is most commonly analyzed with Poisson or negative-binomial regression:

$$\log\{\mathbb{E}(SR_i)\} = \alpha + \mathbf{x}_i^\top \boldsymbol{\beta} \quad (2)$$

This is simple and interpretable. But it rests on four assumptions that are routinely violated in practice:

1. Species are **independent** of each other
2. All species respond **identically** to the environment
3. The species pool is **infinite** (unbounded distribution)
4. The Poisson or NB distribution is **appropriate**

Assumption 1: species independence

Species are not independent; they share environmental responses and co-occur non-randomly.

Assumption 1: species independence

Species are not independent; they share environmental responses and co-occur non-randomly.

If species co-occurrence is positive, species richness has **less** variance than a Poisson. If it is negative, species richness has **more** variance.

Assumption 1: species independence

Species are not independent; they share environmental responses and co-occur non-randomly.

If species co-occurrence is positive, species richness has **less** variance than a Poisson. If it is negative, species richness has **more** variance.

A Poisson or NB regression cannot accommodate this: it will over- or underestimate uncertainty.

Assumption 2: homogeneity of niches

A Poisson or NB distribution for species richness is only justified if all species respond the same way to the environment, i.e., have the same niche.

Assumption 2: homogeneity of niches

A Poisson or NB distribution for species richness is only justified if all species respond the same way to the environment, i.e., have the same niche.

This is ecologically unrealistic: species have different distributions, different ranges, different optimums.

Assumption 2: homogeneity of niches

A Poisson or NB distribution for species richness is only justified if all species respond the same way to the environment, i.e., have the same niche.

This is ecologically unrealistic: species have different distributions, different ranges, different optimums.

The consequence: estimated coefficients from a Poisson richness regression are a **weighted average** over species, with weights depending on the distribution of species' niches, usually not the average ecologists want.

Assumption 3: an infinite species pool

A Poisson distribution is unbounded above: it can predict species richness larger than the number of species in your data.

Assumption 3: an infinite species pool

A Poisson distribution is unbounded above: it can predict species richness larger than the number of species in your data.

This is ecologically impossible. You cannot observe more species than those you have observed.

Negative-binomial adds overdispersion, making the problem worse.

Assumption 3: an infinite species pool

A Poisson distribution is unbounded above: it can predict species richness larger than the number of species in your data.

This is ecologically impossible. You cannot observe more species than those you have observed.

Negative-binomial adds overdispersion, making the problem worse.

Species richness is inherently bounded by the number of species in the observed pool.

The correct distribution: Poisson-Binomial

Under species independence and Bernoulli responses, species richness follows a **Poisson-Binomial distribution**:

$$SR_i \sim \mathcal{PB}(\pi_{i1}, \pi_{i2}, \dots, \pi_{im}) \quad (3)$$

with

$$\mathbb{E}(SR_i) = \sum_{j=1}^m \pi_{ij}, \quad \text{var}(SR_i) = \sum_{j=1}^m \pi_{ij}(1 - \pi_{ij}) \quad (4)$$

The correct distribution: Poisson-Binomial

Under species independence and Bernoulli responses, species richness follows a **Poisson-Binomial distribution**:


$$SR_i \sim \mathcal{PB}(\pi_{i1}, \pi_{i2}, \dots, \pi_{im}) \quad (3)$$

with

$$\mathbb{E}(SR_i) = \sum_{j=1}^m \pi_{ij}, \quad \text{var}(SR_i) = \sum_{j=1}^m \pi_{ij}(1 - \pi_{ij}) \quad (4)$$

- ▶ Mean richness = **sum of species' occurrence probabilities**
- ▶ Variance is bounded (cannot exceed $m/4$)
- ▶ This is the distribution that actually generates species richness
- ▶ The Binomial is a special case (all π_{ij} equal)

Why univariate analysis fails



The Poisson distribution:

- ▶ Puts probability mass on impossible values ($SR > m$)
- ▶ Has the wrong shape (underdispersed here)
- ▶ Assumes the wrong mean-variance relationship

The mismatch worsens when species niches are heterogeneous.

Species richness as an emergent property

We do not model species richness directly; instead, **species richness emerges** from each species' response to the environment:

$$\mathbb{E}(SR_i) = \sum_{j=1}^m \mathbb{E}\{I(y_{ij} > 0)\} = \sum_{j=1}^m p(y_{ij} > 0) \quad (5)$$

Species richness as an emergent property

We do not model species richness directly; instead, **species richness emerges** from each species' response to the environment:

$$\mathbb{E}(SR_i) = \sum_{j=1}^m \mathbb{E}\{I(y_{ij} > 0)\} = \sum_{j=1}^m p(y_{ij} > 0) \quad (5)$$

- ▶ No separate richness model needed
- ▶ Works for any data type: binary, counts, cover, biomass, ordinal
- ▶ Uncertainty propagates naturally from species-level predictions
- ▶ Assumptions are explicit and testable
- ▶ Non-linear richness gradients emerge even from linear species responses

Works for any data type

For any distribution with zero support, the probability of presence can be derived:

Data type	Distribution	$p(y_{ij} > 0)$
Binary	Bernoulli	π_{ij}
Count	Poisson	$1 - e^{-\mu_{ij}}$
Count	Negative-binomial	$1 - (1 + \mu_{ij}/k)^{-k}$
Biomass	Tweedie	$1 - \text{Tw}(0; \mu_{ij}, \phi_j, p)$
Cover	Ordered beta	$1 - \rho_{ij}^0$

Works for any data type

For any distribution with zero support, the probability of presence can be derived:

Data type	Distribution	$p(y_{ij} > 0)$
Binary	Bernoulli	π_{ij}
Count	Poisson	$1 - e^{-\mu_{ij}}$
Count	Negative-binomial	$1 - (1 + \mu_{ij}/k)^{-k}$
Biomass	Tweedie	$1 - \text{Tw}(0; \mu_{ij}, \phi_j, p)$
Cover	Ordered beta	$1 - \rho_{ij}^0$

The only requirement: the distribution must assign probability mass at zero. Gamma and beta distributions do not.

Implementation in gllvm

The predictSR() function computes:

- ▶ Expected species richness $\hat{S}_i = \sum_j \hat{p}_{ij}$
- ▶ Full Poisson-Binomial distribution at each site
- ▶ Prediction intervals; optionally for new predictor values
- ▶ Per-species AUC via goodnessOfFit(): well-calibrated $\hat{p}_{ij}$ give well-calibrated $\hat{S}_i$

```
sr_pred <- predictSR(model, newX = newdata,
                      newLV = matrix(0, ncol = 2, nrow = nrow(
sr_pred$expected      # expected species richness
residuals(sr_pred)    # Dunn-Smyth residuals
goodnessOfFit(model)  # AUC
```

Example: alpine plants in Switzerland

- ▶ Data by D'Amen et al. (2018)
- ▶ Occurrence of 175 species at 840 plots
- ▶ Binary data (presence-absence)
- ▶ Sampled on an elevation gradient in Switzerland

```
dim(Y)
```

```
## [1] 840 175
```


Example: fit a JSDM

```
## gllvm
##      12
## attr(,"autopar")
## [1] TRUE
```

```
model <- gllvm(y = Y[,c(59,(1:ncol(Y))[-59])], X = X_scaled,
               formula = ~ELEV,
               num.lv = 2,
               family = "binomial",
               sd.errors = FALSE)
```

We include elevation as a fixed covariate and two latent variables to capture residual species associations.


Example: a Poisson GLM for species richness




A Poisson GLM with linear elevation predicts a monotone richness curve. Does the multispecies model agree?

Example: LV-elevation confounding

```
obs_sr <- rowSums(Y)
sr_full <- predictSR(model, se.fit = FALSE)
sr_marg <- predictSR(model, se.fit = FALSE, newLV = matrix(0, ncol = 2, nr
```



Example: SR residuals, linear elevation term



Marginal residuals ($LV = 0$) reveal the arch the linear fixed effect misses; conditional residuals are flat because the latent variables absorb the curvature.



Example: adding a quadratic term

```
model_lin <- model
model <- glllvm(y = Y[,c(59,(1:ncol(Y))[-59])], X = X_scaled,
               formula = ~ELEV + I(ELEV^2),
               num.lv = 2,
               family = "binomial",
               sd.errors = TRUE)
```


Species richness peaks at intermediate elevation. Without $ELEV^2$ in the fixed effects, the latent variables absorb the curvature and new-site predictions fail. Adding the quadratic keeps the nonlinearity in the fixed effects where it belongs.

Example: why `plot(model)` does not reveal this

- ▶ `plot(model)` uses **conditional** fitted values: latent variables are estimated per site and adapt to each model
- ▶ A missing quadratic term is absorbed by the latent variables; conditional residuals look clean for both models
- ▶ Marginal predictions ($LV = 0$) reveal the discrepancy




Example: comparing predictions



Question: Both models include a quadratic elevation term. Why do their predictions differ?

Example: residual diagnostics



Poisson-Binomial quantile residuals: no systematic elevation trend remains after adding the quadratic term.

Beyond species richness: the link scale

- ▶ Species richness: known distribution on the response scale (Poisson-Binomial)

Beyond species richness: the link scale

- ▶ Species richness: known distribution on the response scale (Poisson-Binomial)
- ▶ Other diversity quantities: response scale is intractable

Beyond species richness: the link scale

- ▶ Species richness: known distribution on the response scale (Poisson-Binomial)
- ▶ Other diversity quantities: response scale is intractable
- ▶ Like R^2 in mixed models (Nakagawa & Schielzeth 2013), work on the **link scale**: $\text{var}_j(\eta_{ij})$ characterises diversity

Alpha, beta, and gamma diversity

- ▶ **Alpha:** local diversity at a single site; how uneven is the community here?
- ▶ **Gamma:** regional diversity; how much variation is there across the whole community?
- ▶ **Beta:** turnover between sites; how different are two communities?

Alpha, beta, and gamma diversity

- ▶ **Alpha:** local diversity at a single site; how uneven is the community here?
- ▶ **Gamma:** regional diversity; how much variation is there across the whole community?
- ▶ **Beta:** turnover between sites; how different are two communities?

In a model, these follow from the **variance-covariance structure** of η_{ij}

Alpha, beta, and gamma diversity

- ▶ **Alpha:** local diversity at a single site; how uneven is the community here?
- ▶ **Gamma:** regional diversity; how much variation is there across the whole community?
- ▶ **Beta:** turnover between sites; how different are two communities?

In a model, these follow from the **variance-covariance structure** of η_{ij}

The following slides present calculations for a specific model.
 Results are model-dependent.

Dominance

For $\eta_{ij} = \beta_{0j} + \mathbf{x}_i^\top \boldsymbol{\beta}_j$ with $\beta_{0j} \sim \mathcal{N}(0, \sigma_{\beta_0}^2)$ and $\boldsymbol{\beta}_j \sim \mathcal{N}(\mathbf{0}, \Sigma_\beta)$:

$$D_i = \mathbf{x}_i^\top \Sigma_\beta \mathbf{x}_i \quad (6)$$

Dominance

For $\eta_{ij} = \beta_{0j} + \mathbf{x}_i^\top \boldsymbol{\beta}_j$ with $\beta_{0j} \sim \mathcal{N}(0, \sigma_{\beta_0}^2)$ and $\boldsymbol{\beta}_j \sim \mathcal{N}(\mathbf{0}, \Sigma_\beta)$:

$$D_i = \mathbf{x}_i^\top \Sigma_\beta \mathbf{x}_i \quad (6)$$

- ▶ Large D_i : strong environmental filtering; a few well-adapted species dominate
- ▶ Small D_i : weak environmental sorting; more even composition

Gamma diversity

Gamma diversity is the expected within-site variation over sites:

$$\gamma = \sigma_{\beta_0}^2 + \mathbb{E}_i(D_i) \quad (7)$$

Gamma diversity

Gamma diversity is the expected within-site variation over sites:

$$\gamma = \sigma_{\beta_0}^2 + \mathbb{E}_i(D_i) \quad (7)$$

Regional diversity decomposes into a baseline $\sigma_{\beta_0}^2$ (species differ even at the environmental centroid) and average dominance $\mathbb{E}_i(D_i)$ (the contribution of environmental sorting across sites).

Pairwise similarity (beta diversity)

Pairwise similarity between sites i and k is the covariance of the linear predictor across species:

$$\text{cov}_j(\eta_{ij}, \eta_{kj}) = \sigma_{\beta_0}^2 + \mathbf{x}_i^\top \Sigma_\beta \mathbf{x}_k \quad (8)$$

Pairwise similarity (beta diversity)

Pairwise similarity between sites i and k is the covariance of the linear predictor across species:

$$\text{cov}_j(\eta_{ij}, \eta_{kj}) = \sigma_{\beta_0}^2 + \mathbf{x}_i^\top \Sigma_\beta \mathbf{x}_k \quad (8)$$

- ▶ $\sigma_{\beta_0}^2$: baseline similarity shared by all site pairs
- ▶ $\mathbf{x}_i^\top \Sigma_\beta \mathbf{x}_k$: environmental similarity; large when sites share the same environment

Turnover

Turnover between sites i and k follows from the covariance structure:

$$\begin{aligned}
 \text{var}_j(\eta_{ij} - \eta_{kj}) &= \text{var}_j(\eta_{ij}) + \text{var}_j(\eta_{kj}) - 2 \text{cov}_j(\eta_{ij}, \eta_{kj}) \\
 &= \sigma_{\beta_0}^2 + D_i + \sigma_{\beta_0}^2 + D_k - 2 \left(\sigma_{\beta_0}^2 + \mathbf{x}_i^\top \Sigma_\beta \mathbf{x}_k \right) \\
 &= (\mathbf{x}_i - \mathbf{x}_k)^\top \Sigma_\beta (\mathbf{x}_i - \mathbf{x}_k) \quad (9)
 \end{aligned}$$

Turnover

Turnover between sites i and k follows from the covariance structure:

$$\begin{aligned}
 \text{var}_j(\eta_{ij} - \eta_{kj}) &= \text{var}_j(\eta_{ij}) + \text{var}_j(\eta_{kj}) - 2 \text{cov}_j(\eta_{ij}, \eta_{kj}) \\
 &= \sigma_{\beta_0}^2 + D_i + \sigma_{\beta_0}^2 + D_k - 2 \left(\sigma_{\beta_0}^2 + \mathbf{x}_i^\top \Sigma_\beta \mathbf{x}_k \right) \\
 &= (\mathbf{x}_i - \mathbf{x}_k)^\top \Sigma_\beta (\mathbf{x}_i - \mathbf{x}_k) \tag{9}
 \end{aligned}$$

The baseline $\sigma_{\beta_0}^2$ cancels. Turnover has the same quadratic form as D_i , evaluated at the difference $\mathbf{x}_i - \mathbf{x}_k$.

Uniqueness

The average pairwise similarity of site i with all other sites is:

$$\mathbb{E}_k \{ \text{cov}_j(\eta_{ij}, \eta_{kj}) \} = \sigma_{\beta_0}^2 + \mathbf{x}_i^\top \Sigma_\beta \bar{\mathbf{x}} \quad (10)$$

Uniqueness

The average pairwise similarity of site i with all other sites is:

$$\mathbb{E}_k \{ \text{cov}_j(\eta_{ij}, \eta_{kj}) \} = \sigma_{\beta_0}^2 + \mathbf{x}_i^\top \Sigma_\beta \bar{\mathbf{x}} \quad (10)$$

Uniqueness is the average dissimilarity of site i from all other sites:

$$\text{uniqueness}_i = \mathbb{E}_k \{ \text{var}_j(\eta_{ij} - \eta_{kj}) \} = \sigma_{\beta_0}^2 + D_i + \gamma - 2 \left(\sigma_{\beta_0}^2 + \mathbf{x}_i^\top \Sigma_\beta \bar{\mathbf{x}} \right) \quad (11)$$

Uniqueness

The average pairwise similarity of site i with all other sites is:

$$\mathbb{E}_k \{ \text{cov}_j(\eta_{ij}, \eta_{kj}) \} = \sigma_{\beta_0}^2 + \mathbf{x}_i^\top \Sigma_\beta \bar{\mathbf{x}} \quad (10)$$

Uniqueness is the average dissimilarity of site i from all other sites:

$$\text{uniqueness}_i = \mathbb{E}_k \{ \text{var}_j(\eta_{ij} - \eta_{kj}) \} = \sigma_{\beta_0}^2 + D_i + \gamma - 2 \left(\sigma_{\beta_0}^2 + \mathbf{x}_i^\top \Sigma_\beta \bar{\mathbf{x}} \right) \quad (11)$$

Sites at environmental extremes have high uniqueness: their communities differ most from the regional average — potential refugia, range edges, endemic communities.

Example: Swiss birds

Swiss breeding bird atlas: 56 species at 1000+ sites, with temperature (bio_1), precipitation (bio_12), and slope.

Example: Swiss birds

Swiss breeding bird atlas: 56 species at 1000+ sites, with temperature (bio_1), precipitation (bio_12), and slope.

Fit a JSMD with random species-specific slopes, no LVs:

```
model_birds <- gllvm(Y, X, formula = ~(bio_1 + bio_12 + slp | 1),
  family = "binomial", num.lv = 0, beta0com = TRUE)
```

Example: Swiss birds

Swiss breeding bird atlas: 56 species at 1000+ sites, with temperature (bio_1), precipitation (bio_12), and slope.

Fit a JSMD with random species-specific slopes, no LVs:

```
model_birds <- gllvm(Y, X, formula = ~(bio_1 + bio_12 + slp | 1),
  family = "binomial", num.lv = 0, beta0com = TRUE)
```

$\hat{\Sigma}_{\beta}$ is then the estimated covariance of random slopes. All diversity quantities follow from it.

Predicting dominance

```
terra::values(alpha) <- rowSums(x * (x %*% model_birds$params$sigmaB))
```

`model_birds$params$sigmaB` is $\hat{\Sigma}_{\beta}$: the estimated covariance matrix of random slopes. This evaluates $D_i = \mathbf{x}_i^{\top} \hat{\Sigma}_{\beta} \mathbf{x}_i$ at each grid cell: the environmental component of alpha diversity (dominance) due to temperature, precipitation, and slope.

Dominance of Swiss birds

Bird dominance (due to temperature, precipitation and slope)



Extending to LVs

Adding $\mathbf{u}_i^\top$ to the linear predictor introduces residual species associations not captured by environment. Each diversity quantity gains an extra term:

- ▶ **Within-site variation:** $\sigma_{\beta_0}^2 + D_i + \mathbf{u}_i^\top \frac{1}{m} \Lambda^\top \Lambda \mathbf{u}_i$
- ▶ **Turnover:** $\Delta \mathbf{x}^\top \Sigma_\beta \Delta \mathbf{x} + \Delta \mathbf{u}^\top \frac{1}{m} \Lambda^\top \Lambda \Delta \mathbf{u}$
- ▶ **Avg. pairwise similarity:** $\sigma_{\beta_0}^2 + \mathbf{x}_i^\top \Sigma_\beta \bar{\mathbf{x}}$, unchanged (LV scores are centered, $\bar{\mathbf{u}} \approx \mathbf{0}$)


Extending to LVs

Adding $\mathbf{u}_i^\top$ to the linear predictor introduces residual species associations not captured by environment. Each diversity quantity gains an extra term:

- ▶ **Within-site variation:** $\sigma_{\beta_0}^2 + D_i + \mathbf{u}_i^\top \frac{1}{m} \Lambda^\top \Lambda \mathbf{u}_i$
- ▶ **Turnover:** $\Delta \mathbf{x}^\top \Sigma_\beta \Delta \mathbf{x} + \Delta \mathbf{u}^\top \frac{1}{m} \Lambda^\top \Lambda \Delta \mathbf{u}$
- ▶ **Avg. pairwise similarity:** $\sigma_{\beta_0}^2 + \mathbf{x}_i^\top \Sigma_\beta \bar{\mathbf{x}}$, unchanged (LV scores are centered, $\bar{\mathbf{u}} \approx \mathbf{0}$)


LVs primarily inflate uniqueness (through alpha) and add residual turnover. Gamma and uniqueness extend analogously.

Example: pairwise similarity




Pairwise similarity decreases with temperature difference. Pairs at opposite ends of the gradient can be negative: their species compositions are anti-correlated, not merely dissimilar.

Example: pairwise turnover



Environmental turnover increases with temperature difference; the LV component adds residual community dissimilarity not explained by the measured predictors. Averaging pairwise turnover over all partners k gives uniqueness: the per-site expected departure from the rest of the landscape.

Example: uniqueness



Sites at temperature extremes are the most unique: their communities differ most from the regional average. These are the sites that drive beta diversity.

Summary

- ▶ Model the community; diversity follows
- ▶ Univariate species richness analyses make strong and routinely violated assumptions
- ▶ Variance components connect to alpha, beta, and gamma diversity; D_i captures environmental dominance
- ▶ Uniqueness quantifies how much a site's community differs from the regional average

Summary

- ▶ Model the community; diversity follows
- ▶ Univariate species richness analyses make strong and routinely violated assumptions
- ▶ Variance components connect to alpha, beta, and gamma diversity; D_i captures environmental dominance
- ▶ Uniqueness quantifies how much a site's community differs from the regional average

Modelling species richness well requires modelling species well first.