

# Beyond vanilla GLLVMs

Bert van der Veen

Department of Mathematical Sciences, NTNU

## Questions so far?

---



## Some topics covered in this workshop

---

- ▶ Vector GLMs, GLMMs and recap (day 1)
- ▶ Model-based unconstrained ordination: UO, UQO (day 2)
- ▶ Ordination with covariates (day 3)
- ▶ JSDMs: 4th corner model, Phylogenetic random effects (day 3)

## Topics not covered in this workshop

---

- ▶ Joint models with imperfect detection ([Tobler et al. 2019](#), [Doser et al. 2023](#), [Hogg et al. 2021](#))
- ▶ Mixture models for clustering species ([Hui et al. 2015](#), [Hui 2017](#))
- ▶ Spatio-temporal analysis
- ▶ Other methods for analysing community ecological data
- ▶ Much, much more.

## Next steps

---

More response types for unimodal response model?

Spatio-temporal modeling is still a challenge.

- ▶ In space and/or time dimensions are usually large
- ▶ This slows down models considerably
- ▶ Requires additional approximations (e.g, as in INLA, CBFM, or NNGPs as in HMSC)

In `gllvm` we mostly think about modeling ecological processes. However,

- ▶ Data are often not randomly sampled
- ▶ Might contain bias
- ▶ Presence-only from citizen scientists
- ▶ Mixed response types

This warrants further developments.

## Next steps for me

---

I have an article on Phylogenetic random effects in the works.  
However, it is the precursor for a few other developments:

- ▶ Improvements in usability (get methods to a larger audience)
- ▶ Incorporate traits and Phylogeny into ordinations
- ▶ Improvements of robustness (optimisation)

Figure 1: Quadriplot  
*ter Braak et al. (2018)*



## Reduced-rank modeling

---

$$g \left\{ \mathbb{E}(\mathbf{Y}|\mathbf{E}) \right\} = \mathbf{1}\beta^{\top} + \mathbf{E}, \quad \epsilon_i \sim \mathcal{N}(\mathbf{0}, \Sigma) \quad (1)$$

This is a GLMM

## Reduced-rank modeling

---

$$g \left\{ \mathbb{E}(\mathbf{Y}|\mathbf{E}) \right\} = \mathbf{1}\beta^\top + \mathbf{E}, \quad \epsilon_i \sim \mathcal{N}(\mathbf{0}, \mathbf{\Gamma}\mathbf{\Gamma}^\top) \quad (1)$$

This is an unconstrained ordination

## Reduced-rank modeling

$$g\left\{\mathbb{E}(\mathbf{Y}|\mathbf{E})\right\} = \mathbf{1}\beta^\top + \mathbf{E}, \quad \epsilon_i \sim \mathcal{N}(\mathbf{0}, \Gamma\Gamma^\top) \quad (1)$$

This is an unconstrained ordination

$$g\left\{\mathbb{E}(\mathbf{Y}|\mathbf{X})\right\} = \mathbf{1}\beta^\top + \mathbf{XB} \quad (2)$$

This is a GLM

## Reduced-rank modeling

$$g\left\{\mathbb{E}(\mathbf{Y}|\mathbf{E})\right\} = \mathbf{1}\beta^\top + \mathbf{E}, \quad \epsilon_i \sim \mathcal{N}(\mathbf{0}, \Gamma\Gamma^\top) \quad (1)$$

This is an unconstrained ordination

$$g\left\{\mathbb{E}(\mathbf{Y}|\mathbf{X})\right\} = \mathbf{1}\beta^\top + \mathbf{X}\mathbf{B}_{lv}\mathbf{\Gamma}^\top \quad (2)$$

This is a constrained ordination

## Reduced-rank modeling

---

$$g \left\{ \mathbb{E}(\mathbf{Y}|\mathbf{E}) \right\} = \mathbf{1}\beta^\top + \mathbf{Z}\mathbf{\Gamma}^\top, \quad \mathbf{z}_i \sim \mathcal{N}(\mathbf{0}, \mathbf{I}) \quad (1)$$

This is an unconstrained ordination

$$g \left\{ \mathbb{E}(\mathbf{Y}|\mathbf{X}) \right\} = \mathbf{1}\beta^\top + \mathbf{Z}\mathbf{\Gamma}^\top \quad (2)$$

This is a constrained ordination

## 4th corner LVM

---

From yesterday:

$$\eta_{ij} = \beta_{0j} + \mathbf{x}_i^\top (\boldsymbol{\beta}_x + \mathbf{b}_j) + \mathbf{tr}_j^\top \mathbf{B}_{xtr} \mathbf{x}_i \quad (3)$$

Usually we have many traits. And many covariates.  
 So, how can we reduce this down to fewer dimensions?

## Hierarchical ordination

$$g \left\{ \mathbb{E}(\mathbf{Y} | \mathbf{Z}, \mathbf{\Gamma}) \right\} = \mathbf{1}\beta^{\top} + \mathbf{Z}\mathbf{\Sigma}\mathbf{\Gamma}^{\top} \quad (4)$$

Intercept  $\xrightarrow{\hspace{10em}}$   $\mathbf{1}\beta^{\top}$   $\xleftarrow{\hspace{10em}}$  Ordination  $\mathbf{Z}\mathbf{\Sigma}\mathbf{\Gamma}^{\top}$


All effects go into the ordination

**Pros:** fewer parameters, easy visualization

**Cons:** generally more difficult to fit/estimate

## Hierarchical ordination

---





## The “Hierarchical” in ordination

Hierarchically modeling the latent variables:

$$\mathbf{Z} = \mathbf{X}_{lv}\mathbf{B}_{lv} + \mathbf{E} \quad (5)$$

1. Fixed effects —————→
2. LV-level errors —————→
3. (Additional random effects)

*van der Veen et al. (2023)*

## The “Hierarchical” in ordination

Hierarchically modeling the latent variables:

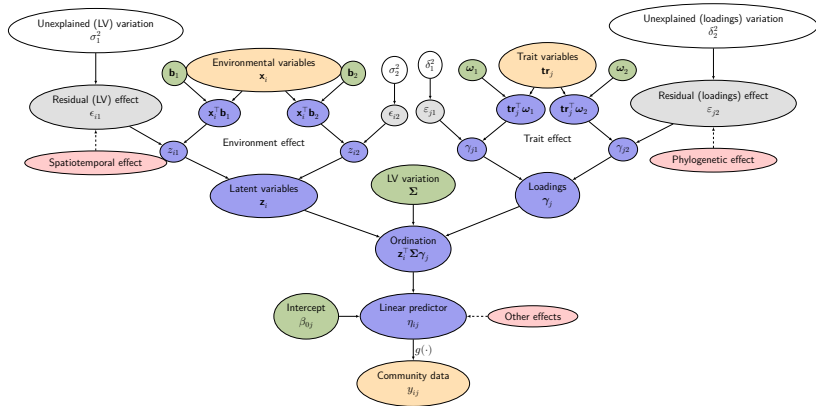
$$\mathbf{Z} = \mathbf{X}_{lv} \mathbf{B}_{lv} + \mathbf{E} \quad (5)$$

1. Fixed effects —————→
2. LV-level errors —————→
3. (Additional random effects)

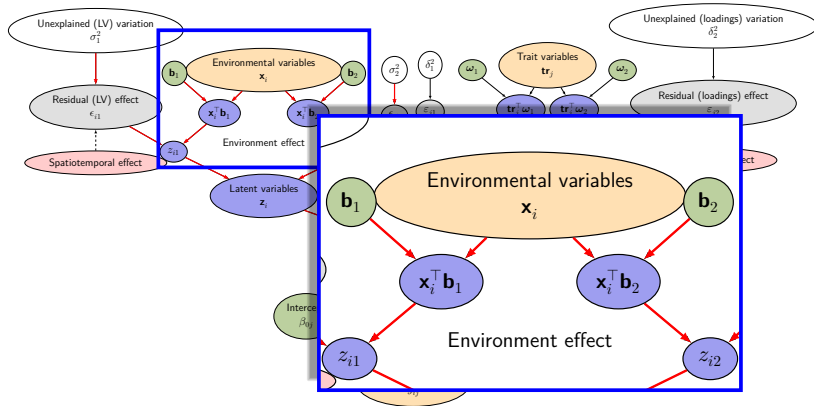
*van der Veen et al. (2023)*

Next: hierarchically modeling the loadings

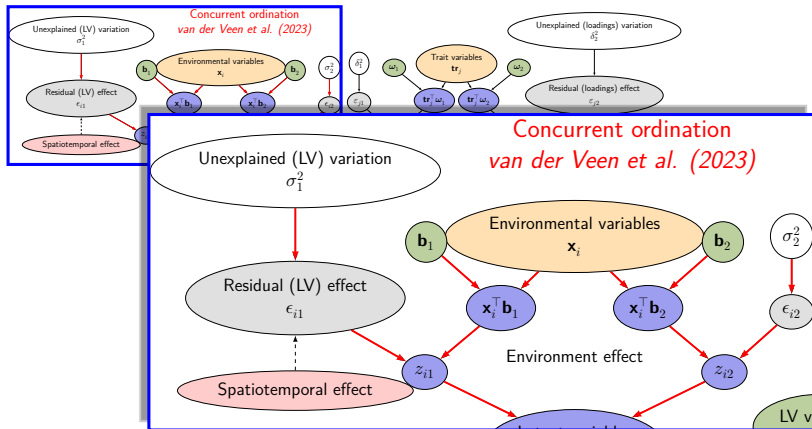
# Hierarchical ordination



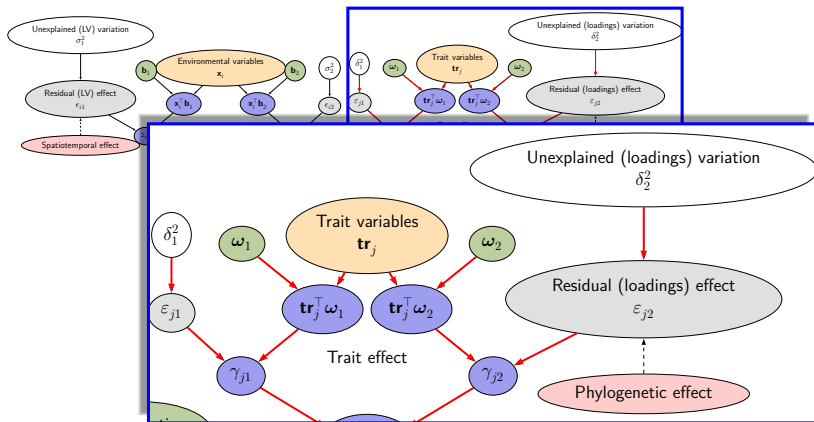
## Hierarchical ordination



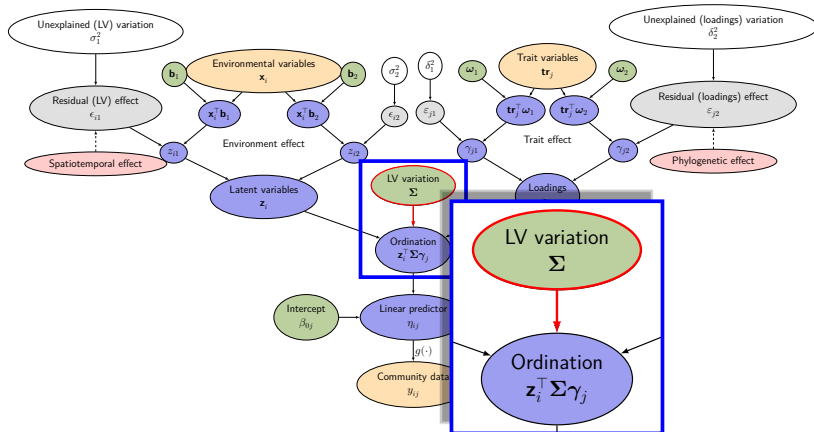
## Hierarchical ordination



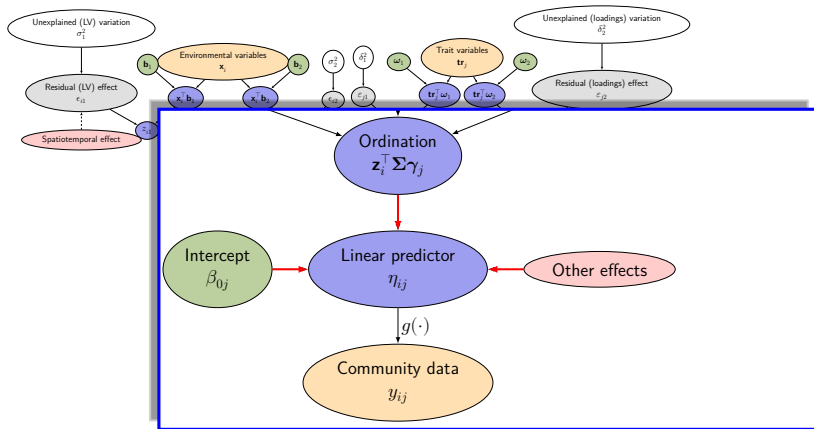
## Hierarchical ordination



## Hierarchical ordination



## Hierarchical ordination





## HO → 4th corner

$$\mathbf{Z} \Sigma \mathbf{\Gamma}^\top =$$

$$\underbrace{\mathbf{XB} \Sigma \mathcal{E}^\top}_{\text{Main effects}} + \overbrace{\mathbf{XB} \Sigma \mathbf{\Omega}^\top \mathbf{TR}^\top}^{\text{4th corner terms}} + \underbrace{\mathbf{E} \Sigma (\mathbf{TR}\mathbf{\Omega} + \mathcal{E})^\top}_{\text{Other stuff}} \quad (6)$$

1. Induces site correlations

2. Induces species correlations

## HO → 4th corner

$$\mathbf{Z} \Sigma \mathbf{\Gamma}^\top =$$

$$\underbrace{\mathbf{XB} \Sigma \mathcal{E}^\top}_{\text{Main effects}} + \overbrace{\mathbf{XB} \Sigma \mathbf{\Omega}^\top \mathbf{TR}^\top}^{\text{4th corner terms}} + \underbrace{\mathbf{E} \Sigma (\mathbf{TR}\mathbf{\Omega} + \mathcal{E})^\top}_{\text{Other stuff}} \quad (6)$$

1. Induces site correlations
2. Induces species correlations

## HO → 4th corner

---

Then we also have:

$$\beta = \mathbf{B} \Sigma \mathcal{E}^{\top} + \mathbf{B} \Sigma \Omega^{\top} \mathbf{T} \mathbf{R}^{\top} \quad (7)$$

I.e., a hierarchical linear regression again of species responses to the environment. The “intercept” term follows from including an intercept column in the trait matrix.

## Example: carabid beetles

---

- ▶ Data from *Ribera et al. (2001)*
  - ▶ Counts
  - ▶ 87 sites
  - ▶ 68 species
  - ▶ Beetles caught with pitfall traps
  - ▶ 26 environmental variables and 27 traits
  - ▶ in Scotland
    - ▶ Management, elevation, soil moisture, pH



## Example: carabid beetles

---



- ▶ Data from *Ribera et al. (2001)*
  - ▶ Counts
  - ▶ 87 sites
  - ▶ 68 species
  - ▶ Beetles caught with pitfall traps
  - ▶ 26 environmental variables and 27 traits
  - ▶ in Scotland
    - ▶ Management, elevation, soil



4887 → 484 parameters (for two dimensions)

## Case study: results\*

- ▶ Reduced-rank approximated terms (e.g., 4th corner)
- ▶ Ordination plots
- ▶ Residual correlation plots



\* fitted with Nimble (de Valpine et al.,

## Hierarchical ordination

---

Read more here:

O'Hara & van der Veen (2022)

## Algorithm-based ordination

---

This workshop focussed strongly on model-based methods for dimension reduction. However, recall that there are three **schools** of ordination:

- 1) "Simple-method"
- 2) "Algorithm-based" (also called "distance-based")
- 3) "Model-based"



## Algorithm-based ordination

---

- ▶ Hierarchical Ordination doubles down on the “model” and ordination angles of these methods
- ▶ JSDMs are in the opposite direction: ordination is an afterthought
- ▶ Algorithm-based methods go the other direction: they double down on the ordination

JSDMs and Algorithm-based ordination both focus on **prediction**

## Algorithm-based ordination

---

A set of methods that we have not covered:

- ▶ Machine learning methods for dimension reduction
- ▶ These can be thought of as similar to NMDS
- ▶ I.e., iterative algorithms that have a geometric interpretation
- ▶ This includes for example: t-SNE, UMAP, Laplacian eigenmaps, isomap, LLE
- ▶ Many, many more
- ▶ No focus (as far as I know) on including covariates

## Algorithm-based ordination

---

NMDS is still very popular, but there have been advances. In all directions, new approaches for co-occurrence analysis are not really adopted in ecology (besides JSDMs).

I will go through a few contemporary algorithm-based ordination methods here.

**Disclaimer: I am not an expert at these methods**

## The paradigm

---

The idea is very different from model-based methods:

- 1) There is no model
- 2) We have a manifold (high dimensional topology)
- 3) Dimension reduction is performed to visualize it in (usually) 2D
- 4) Sites are located the that manifold, and we want to represent their association (no interest in species)

Common in (for example) processing large amounts of data; image processing. Or a biological example: morphometrics.

## Topologies

---

- ▶ In ecology, we do not (yet) have hypotheses about high dimensional topologies
- ▶ Well, kind of. The unimodal model
- ▶ If there is a shape, it is usually consider an “artifact”

Consequently, not many of these methods have made it to ecology.

## Clustering




---

Clustering and ordination are different things. But:

- ▶ Often has focus for algorithm-based ordination
- ▶ A good method supposedly retrieves either, or both, global and local structure
  - ▶ Global: overall patterns, clusters, orderings
  - ▶ Local: relationship of sites that are close together
- ▶ Many of these methods artificially introduces clusters when they do not exist

## Example with handwritten letters


```
data(usps, package = "Rdimtools")  
image(t(matrix(usps$data[4400,], nrow=16)[16:1,]))  
image(t(matrix(usps$data[9900,], nrow=16)[16:1,]))  
image(t(matrix(usps$data[6600,], nrow=16)[16:1,]))
```




# GLLVM

```
HW = t(matrix(usps$data[6600,],nrow=16)[16:1,]) # 6
model<-gllvm::gllvm(HW, num.lv = 2, family = "negative.binomial",
  starting.val="res",quadratic=TRUE, n.init = 1, seed = 2)
image(predict(model,type="response"), main = "Prediction")
image(HW, main = "Data")
```

Prediction



Data





## t-nearest Stochastic Neighbour Embedding ("t-SNE")


---

t-SNE was popular a few years ago. Developed by **van der Maarten (2008)**

- ▶ Mostly because it is very fast
- ▶ Focus is on placing dissimilar sites far away
- ▶ Distances in the data are represented with a normal distribution, in the ordination with a t-distribution
- ▶ Just as in VA it minimizes Kullback-Leibler divergence

## Example with Podani data


```
ord1 <- labdsv::besttsne(vegan::vegdist(PM1), perplexity = 13)  
plot(ord1)
```



Here, “perplexity” is a trade-off between local and global structure.





## Example with the Minchin data

```
ord2 <- labdsv::besttsne(vegan::vegdist(MC), perplexity=15)  
plot(ord2)
```



OK, we can see the lattice!


## Example with the Minchin data



Not better than NMDS or GLLVMs, but definitely not bad.

## Example with handwritten letters


```
ord3 <- labdsv::besttsne(vegan::vegdist(HW),perplexity=5)  
plot(ord3)
```



## Example the dune dataset

Finally a real dataset:

```
ord3 <- labdsv::besttsne(vegan::vegdist(dune), perplexity=6)
plot(ord3)
```



Now we see clusters of sites that might not be real.

## Uniform manifold approximation and projection for dimension reduction (“UMAP”)


---

Is even more recent than t-SNE, developed by **McInnes and Healey (2018)**

- ▶ Supposed to be an improvement over t-SNE
- ▶ Better ability to retain global structure
- ▶ Faster and scales better

## Example with Podani data

```
ord4 <- umap::umap(MC)
plot(ord4$layout, xlab = "UMAP 1", ylab = "UMAP 2")
```




Various R-packages implement UMAP, others have some settings to tune (“epochs”, and nearest neighbours in `uwot`).

This looks terrible (`vegdist` did not improve it much).



## Example with handwritten letters

```
ord5 <- umap::umap(HW)  
plot(ord5$layout)
```



Clearly, it does not do better than t-SNE.

## Locally-linear embedding


---

Introduced by **Roweis and Saul (2000)**

- ▶ Also based on finding a fixed number of nearest neighbours
- ▶ Based on euclidean distance
- ▶ McCune and Grace (2002): it is more similar to PCA than to NMDS

## Example with Podani data

```
ord6 <- Rdimtools::do.lle(as.matrix(PM1), type = c("proportion", 0.2))
plot(ord6$Y, xlab = "LLE1 1", ylab = "LLE 2")
```




It made us a heart. But clearly not a good method.

## Laplacian eigenmaps

Also a method based on eigendecomposition, by Belkin and Niyogi (2003)

```
ord7 <- Rdimtools::do.lapeig(PM1, type=c("proportion", 0.5))
plot(ord7$Y, xlab = "Dimension 1", ylab = "Dimension 2")
```



## isomap


---

Isomap, also known as Principal Curves (Hastie 1984)

- ▶ Has a constrained variant (De'earth 1999)
- ▶ Uses geodesic distances
- ▶ Various parameters to tune: `epsilon`, `k`

## Example with Podani data

```
ord8 <- vegan::isomap(vegan::vegdist(PM1), k = 3)
plot(ord8)
```



## Constrained algorithm-based ordination

---

There are some constrained methods

- ▶ constrained NMDS exists (and not difficult to implement)
- ▶ db-RDA

But general focus is on unconstrained ordination.

## Next generation ordination

---

I have given a brief overview of a few alternative ordination methods.

- ▶ Context is important to see: gllvm performs very well
- ▶ GLLVMs have a wide range of tools for inference, but also prediction
- ▶ They could be adjusted towards topological spaces, but why would we?
- ▶ Dimension reduction is a very broad class of methods
- ▶ Developments are ongoing, and it is important to keep an eye on the competition!

**One thing is clear: some very good methods have been developed in the last decades.**